


# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering



Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>



# Reading

- Szeliski, Chapter 3.1-3.2

# What is an image?




# What is an image?




**We'll focus on these in this class**

**(More on this process later)**



# What is an image?

- A grid (matrix) of intensity values



=

255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255
255	255	127	145	145	175	127	127	95	47	255	255
255	255	74	127	127	127	95	95	95	47	255	255
255	255	255	74	74	74	74	74	74	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)




# What is an image?

- We can think of a (grayscale) image as a **function**,  $f$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$  (or a 2D *signal*):
  - $f(x, y)$  gives the **intensity** at position  $(x, y)$



[snoop](#)



[3D view](#)

- A **digital** image is a discrete (**sampled, quantized**) version of this function

# Image transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (linear filtering)

# Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

What's the next best thing?

# Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function




	7	

Modified image data

# Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



# Cross-correlation

Let  $F$  be the image,  $H$  be the kernel (of size  $2k+1 \times 2k+1$ ), and  $G$  be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

# Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)


$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

- Convolution is **commutative** and **associative**

# Convolution





# Mean filtering


$H$



0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F$

=

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$G$

# Linear filters: examples



Original



0	0	0
0	1	0
0	0	0



Identical image

# Linear filters: examples



Original



0	0	0
1	0	0
0	0	0



Shifted left  
By 1 pixel

# Linear filters: examples




Original



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Blur (with a mean filter)

# Linear filters: examples




Original

$$* \left( \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \right) =$$




**Sharpening filter**  
(accentuates edges)

# Sharpening





before





after

# Smoothing with box filter revisited




# Gaussian Kernel





$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$





# Gaussian filters




$\sigma = 1$  pixel



$\sigma = 5$  pixels




$\sigma = 10$  pixels



$\sigma = 30$  pixels


# Gaussian filter

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian




- Convolving two times with Gaussian kernel of width  $\sigma$  = convolving once with kernel of width  $\sigma\sqrt{2}$

# Sharpening




**before**




**after**

# Sharpening revisited

- What does blurring take away?




Let's add it back:



# Sharpen filter


$$\underset{\substack{\uparrow \\ \text{image}}}{F} + \alpha \left( F - \underbrace{F * H}_{\substack{\text{blurred} \\ \text{image}}} \right)$$

$\uparrow$   
unit impulse  
(identity)




scaled impulse

—



Gaussian

$\approx$



Laplacian of Gaussian

# Sharpen filter





# Convolution in the real world

## Camera shake



Source: Fergus, *et al.* "Removing Camera Shake from a Single Photograph", SIGGRAPH 2006

**Bokeh:** Blur in out-of-focus regions of an image.





Source: <http://lullaby.homepage.dk/diy-camera/bokeh.html>

# Questions?




# Edge detection



- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels


# Origin of Edges



- Edges are caused by a variety of factors

# Characterizing edges

- An edge is a place of rapid change in the image intensity function



# Image derivatives

- How can we differentiate a *digital* image  $F[x,y]$ ?
  - Option 1: reconstruct a continuous image,  $f$ , then compute the derivative
  - Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_x$




$$\frac{\partial f}{\partial y} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_y$

# Image gradient

- The *gradient* of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The *edge strength* is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:


$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?


# Image gradient




$f$



$\frac{\partial f}{\partial x}$




$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$




$\frac{\partial f}{\partial y}$

# Effects of noise




Noisy input image

$$f(x)$$

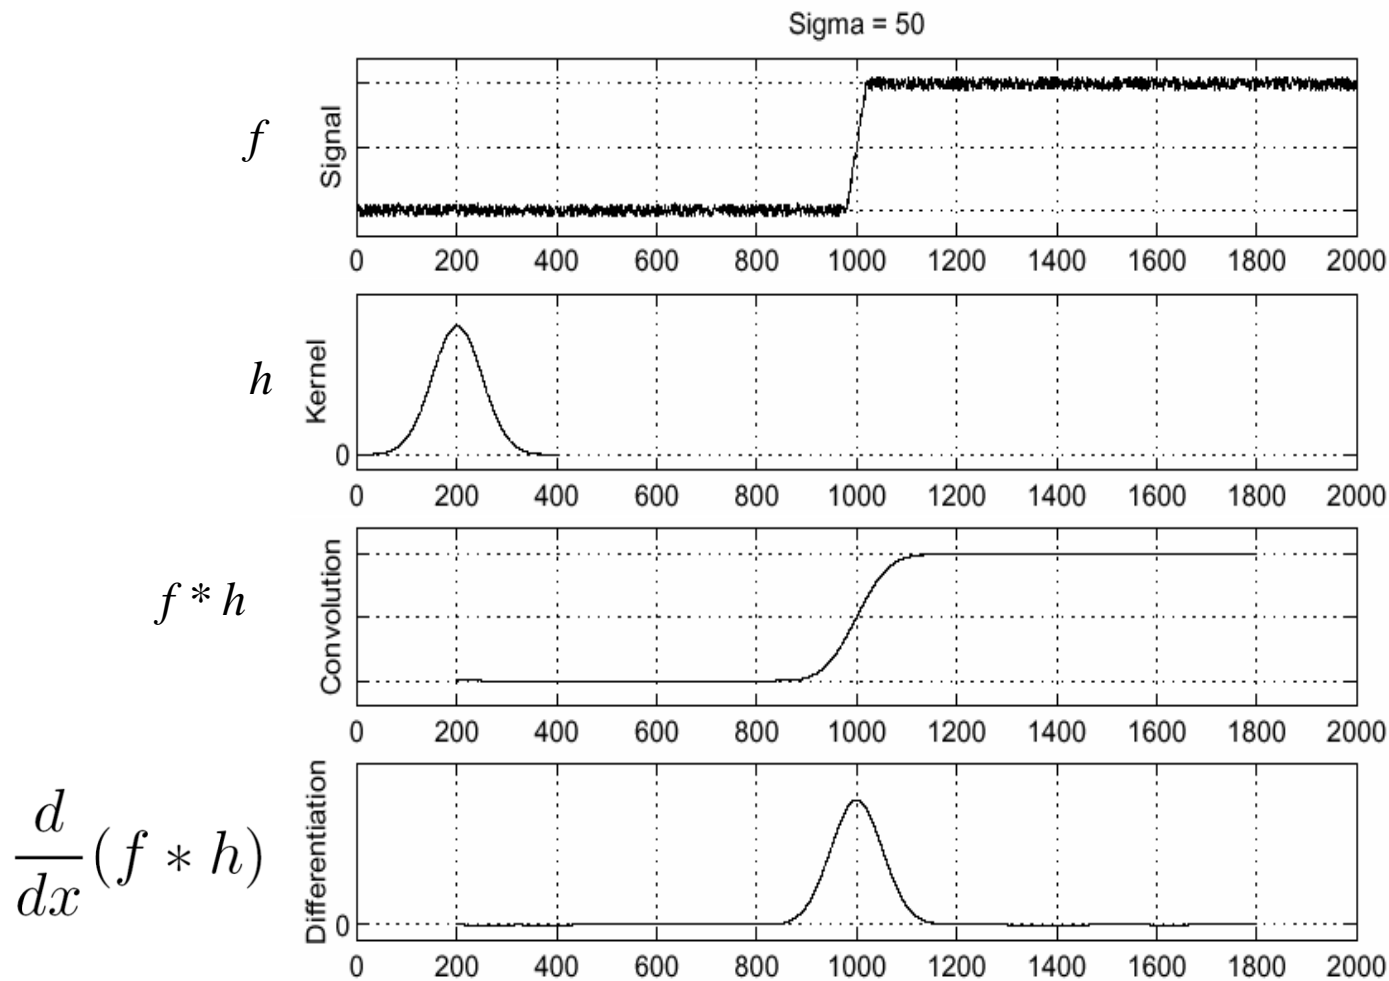


$$\frac{d}{dx}f(x)$$



Where is the edge?

# Solution: smooth first




To find edges, look for peaks in  $\frac{d}{dx}(f * h)$




# Associative property of convolution

- Differentiation is convolution, and convolution is associative:  $\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$
- This saves us one operation:




# 2D edge detection filters



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$




derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$



# Derivative of Gaussian filter





x-direction



y-direction



Side note: How would you compute a directional derivative?



(From vector calculus)


$$\nabla_{\vec{u}} f(\vec{x}) = \nabla f(\vec{x}) \cdot \vec{u}$$

$$\nabla_{\vec{u}} f = ?$$


Directional deriv. is a linear combination of partial derivatives

$$\begin{array}{c} f \\ \begin{array}{ccc} \begin{array}{c} \text{[Radial Pattern]} \end{array} & \begin{array}{c} \text{[Radial Pattern]} \end{array} & \begin{array}{c} \text{[Radial Pattern]} \end{array} \\ \frac{\partial f}{\partial x} \cdot u_x & + & \frac{\partial f}{\partial y} \cdot u_y = \nabla_{\vec{u}} f \end{array}$$

# Derivative of Gaussian filter




x-direction




y-direction


$\cos(\theta)$



$+$   $\sin(\theta)$



$=$



# The Sobel operator

- Common approximation of derivative of Gaussian

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$



$s_x$

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

$s_y$

- The standard defn. of the Sobel operator omits the  $1/8$  term
  - doesn't make a difference for edge detection
  - the  $1/8$  term **is** needed to get the right gradient value

# Sobel operator: example



# Example



- original image (Lena)




# Finding edges



gradient magnitude




# Finding edges



where is the edge?

thresholding

# Non-maximum suppression



- Check if pixel is local maximum along gradient direction
  - requires *interpolating* pixels p and r

# Finding edges



thresholding

# Finding edges



thinning

(non-maximum suppression)



# Canny edge detector

MATLAB: `edge(image, 'canny')`



1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression
4. Linking and thresholding (hysteresis):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them



# Canny edge detector

- Still one of the most widely used edge detectors in computer vision

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

- Depends on several parameters:

$\sigma$  : width of the Gaussian blur

high threshold

low threshold

# Canny edge detector



original



Canny with  $\sigma = 1$




Canny with  $\sigma = 2$

- The choice of  $\sigma$  depends on desired behavior
  - large  $\sigma$  detects “large-scale” edges
  - small  $\sigma$  detects fine edges



# Scale space (Witkin 83)




- Properties of scale space (w/ Gaussian smoothing)
  - edge position may shift with increasing scale ( $\sigma$ )
  - two edges may merge with increasing scale
  - an edge may **not** split into two with increasing scale

# Questions?


- 3-minute break

# Images as vectors


- Very important idea!



2D image



Scanline (1D signal)



Vector

(A 2D,  $n \times m$  image can be represented by a vector of length  $nm$  formed by concatenating the rows)

# Multiplying row and column vectors


$$\begin{bmatrix} 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = ?$$

# Filtering as matrix multiplication


$$\begin{bmatrix}
 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 \\
 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 \\
 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 & 0.2 \\
 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2 & 0.2 \\
 0.2 & 0.2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0.2 & 0.2 & 0.2
 \end{bmatrix}
 \begin{bmatrix}
 0 \\
 0 \\
 0 \\
 0 \\
 0 \\
 0 \\
 0 \\
 0 \\
 0 \\
 1 \\
 1 \\
 1 \\
 1 \\
 1 \\
 1 \\
 1 \\
 1 \\
 1
 \end{bmatrix}$$

What kind of filter is this?


# Filtering as matrix multiplication



# Another matrix transformation




1D Discrete cosine transform (DCT) basis




2D DCT basis

# Another matrix transformation



1D Discrete cosine transform (DCT) basis



2D DCT basis



+

6.192 x